

AI Factory Benchmarking Framework

Automated Performance Analysis for HPC-Based AI Services

Mario Capodanno Giuseppe Galardi Can Beydogan Thies Weel

January 2026

EUMASTER4HPC 2025–2026

Supervisor: Dr. Farouk Mansouri, LuxProvide

Outline

Introduction

System Architecture

Recipe-Driven Benchmarking

Automated Analysis

Real-Time Monitoring

Experimental Results

User Interface

Conclusion

Introduction

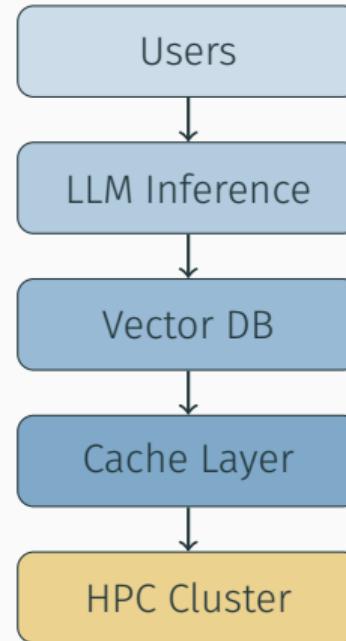
The Challenge: AI Factories on HPC

Modern AI Infrastructure

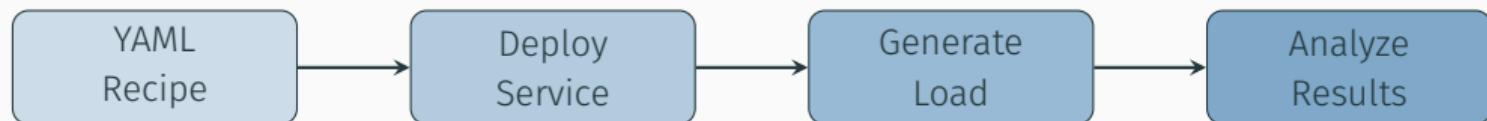
- LLM inference at scale
- Vector databases for RAG
- High-throughput caching
- GPU-accelerated workloads

Key Questions

- Max sustainable concurrency?
- Where are the bottlenecks?
- How do configs affect performance?



Automated Benchmarking Framework



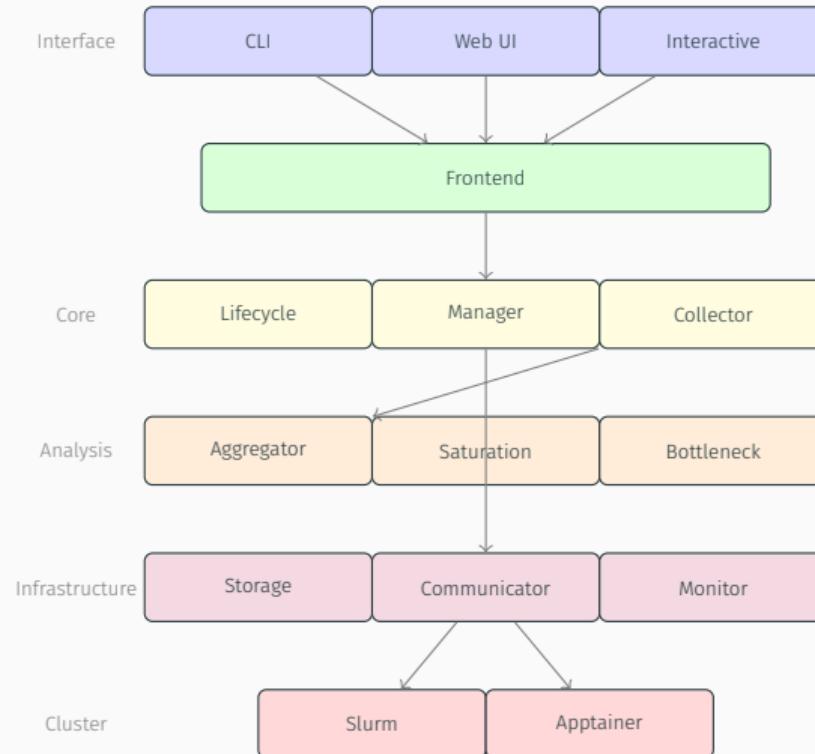
✓ Recipe-driven

✓ Fully automated

✓ Reproducible

System Architecture

Component Overview



Supported Services

Category	Service	Use Case
Inference	vLLM	High-perf LLM serving
	Ollama	Lightweight LLM
Database	PostgreSQL	OLTP workloads
	Redis	Caching
	MinIO	Object storage
Vector DB	ChromaDB	Embedding search
	Qdrant	RAG pipelines

7 services 3 categories Automatic configuration

Recipe-Driven Benchmarking

Recipe Example

recipe_vllm.yaml

```
configuration:
  target: meluxina

service:
  type: vllm
  partition: gpu
  num_gpus: 1
  settings:
    model: facebook/opt-125m

client:
  type: vllm_smoke
  partition: cpu

benchmarks:
  num_clients: 4
```

Benefits

- Self-documenting
- Version-controllable
- Reproducible
- Schema-validated

One Command

```
python frontend.py recipe.yaml
```

Automated Analysis

Saturation Detection

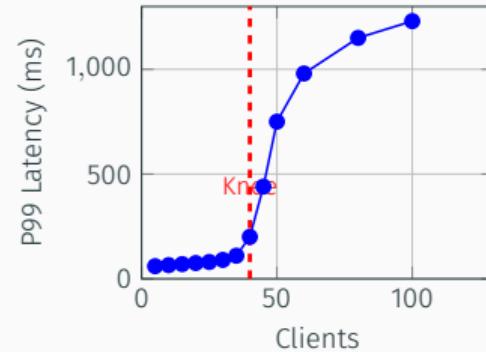
Maximum Curvature Method

Find the “knee” where latency grows superlinearly:

$$\kappa(x) = \frac{|y''(x)|}{(1 + y'(x)^2)^{3/2}}$$

Identifies:

- Latency knee point
- Throughput saturation
- SLO compliance limit



Bottleneck Attribution

Type	Indicators
 GPU-bound	GPU util > 80%, stable CPU, rising TTFT
 CPU-bound	High CPU time, low GPU, stable memory
 Memory-bound	High RSS, OOM errors, latency spikes
 Queueing	Throughput plateau, P99 explosion

Each classification includes evidence and actionable recommendations

Real-Time Monitoring

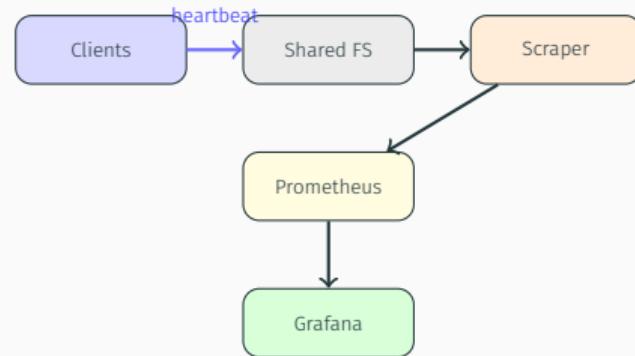
Prometheus + Grafana Integration

Monitoring Stack

- Prometheus for metrics collection
- Grafana for visualization
- Pre-configured dashboards

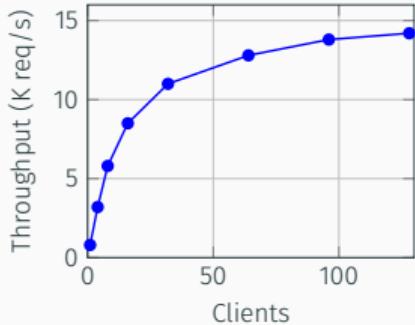
Heartbeat Strategy

- Filesystem-based client tracking
- Works across distributed nodes
- Uses shared Lustre filesystem



Experimental Results

Case Study: Redis Scaling



Findings

- Near-linear scaling to 64 clients
- Saturation at ~96 clients
- Peak throughput: 14.2K req/s
- Bottleneck: CPU processing

Payload Sensitivity

- Small values (<1KB): linear
- Large values (>16KB): earlier saturation due to memory bandwidth

User Interface

CLI Workflow

```
# Run benchmark
python src/frontend.py recipe.yaml

# Monitor
python src/frontend.py --watch BM-20260112-001

# Generate report
python src/frontend.py --report BM-20260112-001

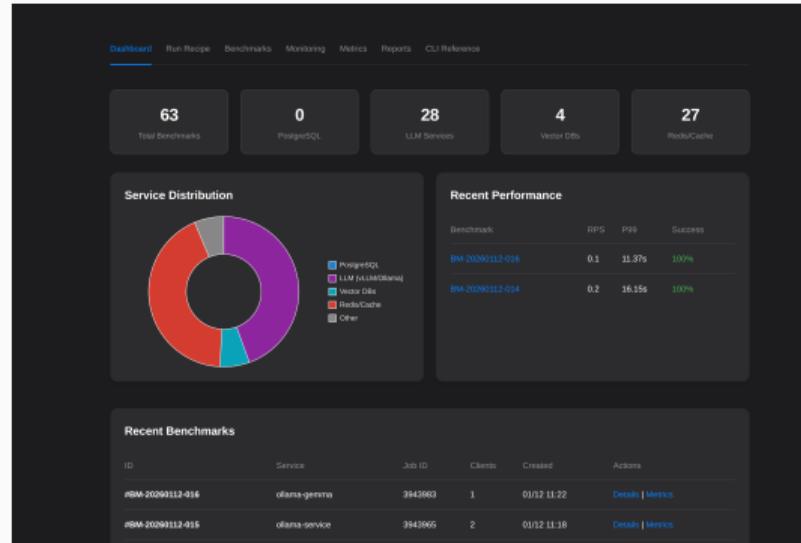
# Compare runs
python src/frontend.py --compare BM-001 BM-002

# Launch web UI
python src/frontend.py --web
```

Web Interface

Features

- 📈 Dashboard
- 🔍 Deploy benchmarks
- 📈 Real-time status
- 📄 Log browser
- 📊 Metrics visualization
- 📄 Report generation



Conclusion

Summary

Key Contributions

- Automated saturation detection
- Bottleneck attribution
- Real time monitoring

Target Platform

- Slurm job scheduling
- Apptainer containers
- GPU and CPU partitions

Future Work

- Multi-node client deployment for higher load
- ML-based anomaly detection
- Integration with PBS, LSF schedulers
- Power-aware efficiency metrics (tokens/watt)

Thank you!



Team1_EUMASTER4HPC2526

Backup: Artifact Structure

```
results/<benchmark_id>/  
  run.json          # Complete metadata  
  requests.jsonl    # Per-request timing  
  summary.json      # Aggregated metrics  
  logs/             # Service and client logs  
  
reports/<benchmark_id>/  
  report.md         # Human-readable analysis  
  report.json       # Machine-readable summary  
  plots/            # Visualization PNGs
```

Backup: Regression Detection

Metric	Threshold
P99 Latency increase	> 10%
Throughput decrease	> 10%
Success rate decrease	> 1%
Error count increase	Any

Configurable thresholds for CI/CD integration